
javautil-dblogging Documentation

Release 19.6.0

Jim Schmidt

October 27, 2019

CONTENTS

Application instrumentation is essential to performance monitoring, unfortunately this is often performed by throwing in some logging statements using a java logging framework or wrapper, such as slf,

This however fails to capture any information essential to end-to-end monitoring as it omits what is generally the biggest source of latency, the relational database. What statements are being executed, how long do they take, which statements take up the bulk of the resources in aggregate (an under-performing statement invoked thousands of times an hour is not uncommon).

Oracle provides, at great expense, the ASH subsystem and even that does not associated the sql statements to the application code.

This utility provides a simple Application Program Interface to allow you to record performance information in a simple, low overhead fashion from any java program or any program that allows pl/sql calls.

Thus a few judicious additions to an Oracle Form or a batch job can provide the foundation of information necessary to establish where database resources are being consumed.

Resolution of these matters is a different matter and may involve gathering statistics, altering execution plans, creating or eliminating indexes or rewriing SQL.

SGA paramaters may have to be changed.

DOCUMENT ORGANIZATION

We will introduce the types of logging, show examples of how to instrument your code and what type of output is generated.

Then we will guide you through the installation and options.

OVERVIEW OF LOGGING TYPES

2.1 Flat file logging

Writes log messages to a text file on the database server.

Includes:

- log_seq_nbr
- log_level
- job_log_id
- log_msg_id
- line_number
- current_timestamp, 'YYYY-MM-DDTHH24:MI:SSXFF'
- log_msg
- caller_name optional
- call_stack optional

This will allow you to easily write log messages from pl/sql in stored procedures, functions, packages and with most of the flexibility and ease of using a java logging framework.

2.1.1 Logging to flat files

Overview

Logs messages using utl_file to a directory on the database server specified.

First the database directory is created and oracle is granted permission to read and write it, then the ddl “create directory....” and “grant read, write on directory...”

Examples

log_to_file_only.proc.sr.sql

Input

```
set serveroutput on
set echo on
create or replace procedure log_to_file_only is
    long_msg clob := 'this is an absurdly long message, ' ||
        ' interesting stuff to say so I will just write meaningless ' ||
        ' stuff for a little while. ' ||
        ' The quick brown fox jumped over the lazy dog. ';

    my_log_file_name varchar(4096);
begin
    my_log_file_name := pllogger.open_log_file('log_to_file_only.text');
    pllogger.set_filter_level(9); -- all messages should go to log file
    pllogger.info('anonymous', $$PLSQL_LINE, 'begin loop');
    pllogger.info($$PLSQL_UNIT, $$PLSQL_LINE, long_msg);
    for i in 1..3
    loop
        pllogger.fine($$PLSQL_UNIT, $$PLSQL_LINE, 'i is ' || to_char(i));
    end loop;
    pllogger.close_log_file();
exception when others then
    -- a severe condition is not necessarily fatal
    pllogger.severe($$PLSQL_UNIT, $$PLSQL_LINE, sqlerrm);
    pllogger.close_log_file();
    raise;
end;
/
show errors

exec log_to_file_only();
```

Output

```
"log_level","job_log_id","job_msg_id","line_number","timestamp","log_msg","caller_name","call_stack"
4,,,17,"2019-10-26T17:19:52.885607","begin loop","anonymous",""
4,,,18,"2019-10-26T17:19:52.886020","this is an absurdly long message, exceeding the length of the l
7,,,22,"2019-10-26T17:19:52.886197","i is 1","LOG_TO_FILE_ONLY",""
7,,,22,"2019-10-26T17:19:52.886357","i is 2","LOG_TO_FILE_ONLY",""
7,,,22,"2019-10-26T17:19:52.886502","i is 3","LOG_TO_FILE_ONLY",""
```

2.2 Database Logging

Record jobs and their steps, how long each step took to execute and optionally extremely detailed information about every database operation as an oracle trace file may be parsed and stored in the log repository.

The log repository may be on the same oracle database server, even the same schema using the same connection as it uses autonomous transactions, or in postgresql or h2.

2.2.1 Database Logging

included file

job_tables

In the interest of expediency we have a quick listing of the job tables.

```
SQL> describe job_log
```

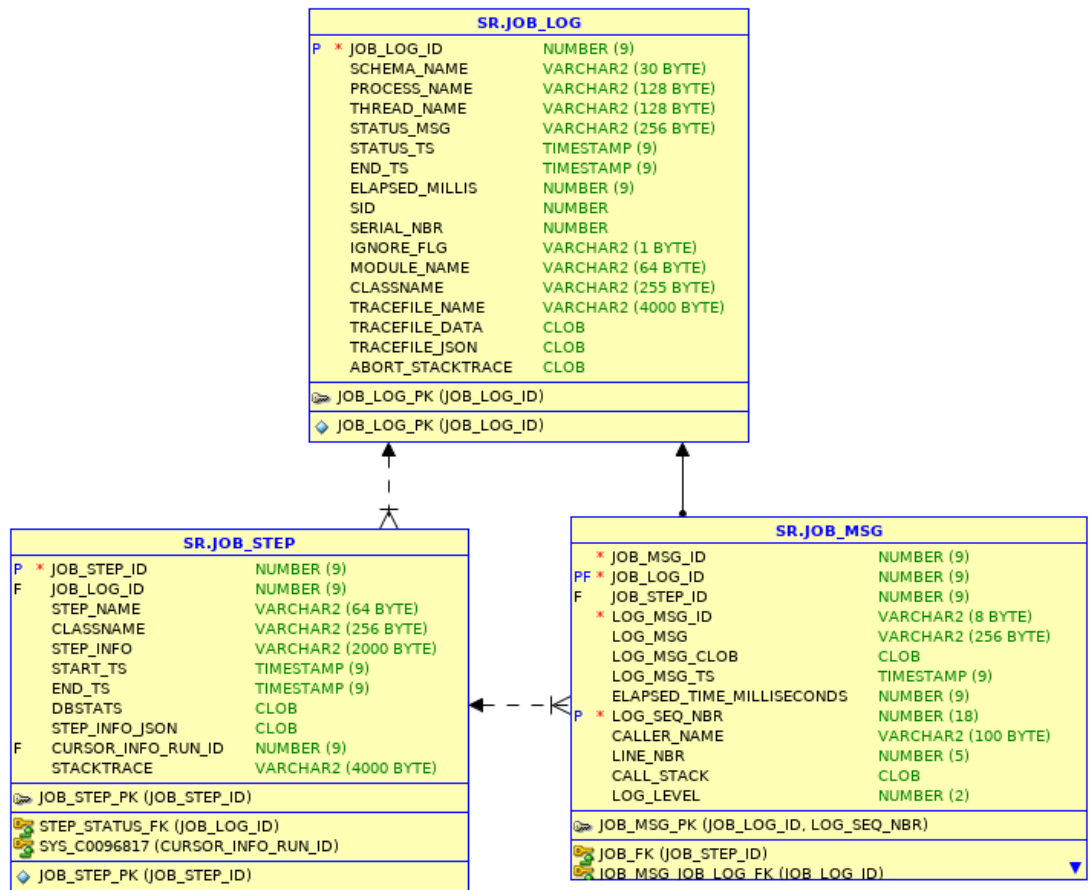
Name	Null?	Type
JOB_LOG_ID	NOT NULL	NUMBER(9)
SCHEMA_NAME		VARCHAR2(30)
PROCESS_NAME		VARCHAR2(128)
THREAD_NAME		VARCHAR2(128)
STATUS_MSG		VARCHAR2(256)
STATUS_TS		TIMESTAMP(9)
END_TS		TIMESTAMP(9)
ELAPSED_MILLIS		NUMBER(9)
SID		NUMBER
SERIAL_NBR		NUMBER
IGNORE_FLG		VARCHAR2(1)
MODULE_NAME		VARCHAR2(64)
CLASSNAME		VARCHAR2(255)
TRACEFILE_NAME		VARCHAR2(4000)
TRACEFILE_DATA		CLOB
TRACEFILE_JSON		CLOB
ABORT_STACKTRACE		CLOB

```
SQL> describe job_step
```

Name	Null?	Type
JOB_STEP_ID	NOT NULL	NUMBER(9)
JOB_LOG_ID		NUMBER(9)
STEP_NAME		VARCHAR2(64)
CLASSNAME		VARCHAR2(256)
STEP_INFO		VARCHAR2(2000)
START_TS		TIMESTAMP(9)
END_TS		TIMESTAMP(9)
DBSTATS		CLOB
STEP_INFO_JSON		CLOB
CURSOR_INFO_RUN_ID		NUMBER(9)
STACKTRACE		VARCHAR2(4000)

```
SQL> describe job_msg;
```

Name	Null?	Type
JOB_MSG_ID		NUMBER(9)
JOB_LOG_ID	NOT NULL	NUMBER(9)
LOG_MSG_ID		VARCHAR2(8)
LOG_MSG		VARCHAR2(256)
LOG_MSG_CLOB		CLOB
LOG_MSG_TS		TIMESTAMP(9)
ELAPSED_TIME_MILLISECONDS		NUMBER(9)
LOG_SEQ_NBR	NOT NULL	NUMBER(18)
CALLER_NAME		VARCHAR2(100)
LINE_NBR		NUMBER(5)
CALL_STACK		CLOB
LOG_LEVEL		NUMBER(2)



Entity Relationship Diagram

TODO run the python with the comments

Each job has one job_log entry and one or more job_steps.

Job steps may have associated log messages.

dblogger_install_tables

Creates the job and job step tables and views

sequences cursor_info_run_id_seq; cursor_info_id_seq; job_log_id_seq; job_msg_id_seq; job_step_id_seq;

tables

- cursor_explain_plan
- cursor_sql_text
- cursor_info_run
- cursor_info
- cursor_stat
- job_log
- job_msg
- job_step

views

- cursor_info_vw
- job_step_vw
- job_log_vw

Job Logging

Logging information may be written to a text file, stored in a database and written to the oracle trace file.

Steps start job logging.

```
public long sampleUsage(Dblogger dblogger, Connection appConnection) throws SqlSplitterException, Ex
    dblogger.prepareConnection();
    final String processName = "Process Name";
    // Start the job

    final long logJobId = dblogger.startJobLogging(processName, getClass().getName(), null, null, 4);
    dblogger.setModule("SplitLoggerTest", "simple example");
    dblogger.setAction("Some work");
    dblogger.insertStep("Full join", "Meaningless busy work", getClass().getName());
    ConnectionUtil.exhaustQuery(appConnection, "select * from user_tab_columns, user_tables where row
```

```
    dblogger.setAction("Another set of work");
    ConnectionUtil.exhaustQuery(appConnection, "select count(*) from all_tab_columns");
    // End the job
    dblogger.endJob();
    return logJobId;
}
```

job logging persistence has a bit of indirection

Installation

Repositories

RDBMS persistence support is provided for Oracle, H2 and postgresql

H2 is a lightweight database and may be used to eliminate the need for support of another Oracle Database.

Postgresql is a high end database that requires minimal installation and administration.

You should probable not compound your problem with yet another Oracle install, but if your DBA will allow you a schema in your database for logging, you don't have to learn anything else.

The Oracle database could be the same instance as the application being monitored, but this may raise some objections to the application DBA.

Oracle logging repository

If the logging data is to be persisted in Oracle, the tables must be created and some packages created.

1. job_log
2. job_msg
3. job step

The granularity of job step is left to the invoker.

As the overhead is very low, there is no reason to be parsimonious

with identification, it's a simple one line call in the user app.

These records can be reviewed for job success or failure and form a historical basis of time elapsed by job and step.

This may be used as a starting pointing in locating "what processes are using the time?"

Additionally they constitute a base performance metric from which runtime degradation or periodic anomalous runs may be identified.

Data is committed by calls from java to the package logger, provided here.

The package utilizes autonomous commits and hence may be safely called using the same connection as the application.

2.3 Oracle trace information

The third type of logging is an extension of database logging and stores oracle trace information a relational database.

Oracle tracing is turned on and the trace files parsed and aggregated and stored in tables associated with the various job steps.

- oracle
- h2
- postgresql

DATABASE LOGS

We will start with an example program and show what is logged.

3.1 Java Example

```
package org.javautil.dblogging;

import java.sql.Connection;
import java.sql.SQLException;

import org.javautil.core.sql.Binds;
import org.javautil.core.sql.ConnectionUtil;
import org.javautil.core.sql.SqlStatement;
import org.javautil.dblogging.logger.Dblogger;
import org.javautil.util.NameValue;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

public class DbloggerForOracleExample {

    private Dblogger dblogger;
    private Connection connection;
    private String processName;
    private boolean testAbort = false;
    private int traceLevel;
    private final Logger logger = LoggerFactory.getLogger(getClass());

    public DbloggerForOracleExample(Connection connection, Dblogger dblogger, String processName, boolean testAbort, int traceLevel) {
        this.connection = connection;
        this.dblogger = dblogger;
        this.processName = processName;
        this.testAbort = testAbort;
        this.traceLevel = traceLevel;
    }

    public long process() throws SQLException {
        dblogger.prepareConnection();
        long id = 0;

        try {
```

```
        id = dblogger.startJobLogging(processName, getClass().getName(), "ExampleLogging", null,
        logger.debug("=====jobId: {}", id);
        limitedFullJoin();
        fullJoin();
        userTablesCount();
        if (testAbort) {
            int x = 1 / 0;
        }
        logger.debug("calling endJob");
        dblogger.endJob();
    } catch (Exception e) {
        logger.error(e.getMessage());
        e.printStackTrace();
        dblogger.abortJob(e);
        throw e;
    }
    return id;
}

/**
 * This will set the v$session.action
 */
private void limitedFullJoin() throws SQLException {
    logger.debug("limitedFullJoin =====");
    dblogger.setAction("actionNoStep");
    ConnectionUtil.exhaustQuery(connection, "select * from user_tab_columns, user_tables where r
    dblogger.setAction(null); // no longer performing that action, so clear
    logger.debug("limitedFullJoin complete =====");
}

private void fullJoin() throws SQLException {
    logger.debug("fullJoinBegins =====");
    // TODO insertStep should set the action
    dblogger.insertStep("fullJoin", "fullJoin", getClass().getName());
    ConnectionUtil.exhaustQuery(connection, "select * from user_tab_columns, user_tables");
    dblogger.finishStep();
    logger.debug("fullJoin complete =====");
}

private void userTablesCount() throws SQLException {
    dblogger.insertStep("count full", "userTablesCount", getClass().getName());
    ConnectionUtil.exhaustQuery(connection, "select count(*) dracula from user_tables");
    dblogger.finishStep();
    // TODO support implicit finish step
}

NameValue getJobLog(Connection connection, long id) throws SQLException {
    final String sql = "select * from job_log " + "where job_log_id = :job_stat_id";
    final SqlStatement ss = new SqlStatement(connection, sql);
    Binds binds = new Binds();
    binds.put("job_stat_id", id);
    final NameValue retval = ss.getNameValue();
    ss.close();
    return retval;
}
}
```

3.2 Analyzing the logs

Separate utilities are used to analyze the logs. A very useful tool is `javautil-condition-identification`.

Did any job abort?

What is the trend on elapsed times?

How do elapsed times vary based on time of day?

Getting deeper, with trace information one can drill down to the details, we will cover that later.

TRACEFILE GENERATION AND PERSISTENCE

This utility provides the information to the Oracle Performance specialist to identify the root cause of the problem, how to repair is another speciality.

INSTALLATION OF DATABASE ARTIFACTS FOR ORACLE

These files may be found under *src/main/resources/ddl/oracle*

The script that runs them all is *install.sql*

```
set echo on
@prepare_connection.sql
@my_session_info.sql
@dblogger_uninstall.sr.sql
@logger_message_formatter.plsql.sr.sql
@dblogger_install_tables.sr.sql
@dblogger_install.pks.sr.sql
@dblogger_install.pkb.sr.sql
@logger_persistence.pks.sr.sql
@logger_persistence.pkb.sr.sql
```

5.1 prepare_connection

prepare-connection provides one procedure.

This will call `dbms_session.clear_context` for each context variable,

restoring the context for a connection returned from a `connection_pool` to the state the of an initially opened connection.

Connection pools do not generally clear this information out as it is Oracle specific.

5.2 my_session_info.sql

creates the view *my_session_info* to allow the connected user to obtain the `v$session` record for the current connection.

5.3 logger_message_formatter

Provides the *logger_message_formatter* function, which creates a single string from all of the logging parameters and makes a call to `dbms_output.put_line` and then returns the formatted message.

5.4 dblogger_intall_tables

Creates the job and job step tables and views

cursor_info_run_id_seq; cursor_info_id_seq; job_log_id_seq; job_msg_id_seq; job_step_id_seq;

- cursor_explain_plan
- cursor_sql_text
- cursor_info_run
- cursor_info
- cursor_stat
- job_log
- job_msg
- job_step
- cursor_info_vw
- job_step_vw
- job_log_vw

5.5 Job Logging

Logging information may be written to a text file, stored in a database and written to the oracle trace file.

5.5.1 Steps start job logging.

```
public long sampleUsage(Dblogger dblogger, Connection appConnection) throws SqlSplitterException, Ex
    dblogger.prepareConnection();
    final String processName = "Process Name";
    // Start the job

    final long logJobId = dblogger.startJobLogging(processName, getClass().getName(), null, null, 4);
    dblogger.setModule("SplitLoggerTest", "simple example");
    dblogger.setAction("Some work");
    dblogger.insertStep("Full join", "Meaningless busy work", getClass().getName());
    ConnectionUtil.exhaustQuery(appConnection, "select * from user_tab_columns, user_tables where row

    dblogger.setAction("Another set of work");
    ConnectionUtil.exhaustQuery(appConnection, "select count(*) from all_tab_columns");
    // End the job
    dblogger.endJob();
    return logJobId;
}
```

job logging persistence has a bit of indirection

5.6 Installation

5.6.1 Repositories

RDBMS persistence support is provided for Oracle, H2 and postgresql

H2 is a lightweight database and may be used to eliminate the need for support of another Oracle Database.

Postgresql is a high end database that requires minimal installation and administration.

You should probable not compound your problem with yet another Oracle install, but if your DBA will allow you a schema in your database for logging, you don't have to learn anything else.

The Oracle database could be the same instance as the application being monitored, but this may raise some objections to the application DBA.

5.6.2 Oracle logging repository

If the logging data is to be persisted in Oracle, the tables must be created and some packages created.

Job log tables

1. job_log
2. job_msg
3. job step

The granularity of job step is left to the invoker.

As the overhead is very low, there is no reason to be parsimonious

with identification, it's a simple one line call in the user app.

These records can be reviewed for job success or failure and form a historical basis of time elapsed.

This may be used as a starting pointing in locating "what processes are using the time?"

Additionally they constitute a base performance metric from which runtime degradation or periodic anomalous runs may be identified.

Data is committed by calls from java to the package logger, provided here.

The package creates autonomous commits and hence may be safely called using the same connection as the application.

5.6.3 logging package

The logger package provides the following:

These primarily set information in the SGA and enable oracle session tracing.

begin_java_java

change v\$session information

```
procedure prepare_connection;  
set_module set action
```

5.7 Trace Repository

- cursor_explain_plan
- cursor_sql_text
- cursor_info_run
- cursor_info
- cursor_stat

5.8 logger_persistence package

The logger persistence package provides an API for writing to various tables using autonomous transactions.

```
procedure save_job_log (  
    p_job_log_id    in number,  
    p_schema_name   in varchar2,  
    p_process_name  in varchar2,  
    p_classname     in varchar2,  
    p_module_name   in varchar2,  
    p_status_msg    in varchar2,  
    p_thread_name   in varchar2,  
    p_trace_level   in pls_integer default logger.G_INFO,  
    p_tracefile_name in varchar2,  
    p_sid           in pls_integer  
);
```

The source of work is indentifiable down to the java thread.

```
function save_job_step (  
    p_job_log_id in pls_integer,  
    p_step_name  in varchar,  
    p_step_info  in varchar,  
    p_classname  in varchar,  
    p_start_ts   in timestamp,  
    p_stacktrace in varchar  
  
) return number;  
  
procedure finish_step  
  
procedure end_job(p_elapsed_milliseconds in pls_integer)  
  
procedure abort_job(p_elapsed_milliseconds in pls_integer,p_stacktrace in varchar);
```

5.8.1 Install Oracle JDBC

See [this post](#) to use Oracle JDBC properly. Or, you could download the JAR file, and then execute this command:

```
TODO the script to locate mvn install:install-file -DgroupId=com.oracle  
-DartifactId=oracle-jdbc8 -Dversion=12c -Dpackaging=jar -Dfile=<THE_JDBC_JAR_LOCATION>
```

Notations in job .sql script used by sqlrunner.

SECURITY IN PRODUCTION

USER PRIVILIGES

PERFORMANCE

TOOLS AND CONCEPTS

User should be familiar with v\$session view, tkprof command line utility

CONNECTION POOLS

10.1 After Getting a connection

10.1.1 Contexts

If a session is being used as part of a connection pool and the state of its contexts are not reinitialized, this can lead to unexpected behavior.

10.1.2 Packages

Sessions have the ability to alter package state by amending the values of package variables. If a session is being used as part of a connection pool and the state of its packages are not reinitialized, this can lead to unexpected behavior. To solve this, Oracle provides the `dbms_session.reset_package` procedure.

The `dblogging` provided procedure clears all context variables and resets package state.

Connections must be reset immediately after being obtained from a connection pool

In `src/main/resources/ddl/oracle/prepare_connection`

10.1.3 Convenience Procedure

```
CREATE OR REPLACE PROCEDURE prepare_connection
AS
    context_info DBMS_SESSION.AppCtxTabTyp;
    info_count PLS_INTEGER;
    indx PLS_INTEGER;
BEGIN
    DBMS_SESSION.LIST_CONTEXT ( context_info, info_count);
    indx := context_info.FIRST;
    LOOP
        EXIT WHEN indx IS NULL;
        DBMS_SESSION.CLEAR_CONTEXT (
            context_info(indx).namespace,
            context_info(indx).attribute,
            null
        );
        indx := context_info.NEXT (indx);
    END LOOP;
    DBMS_SESSION.RESET_PACKAGE;
END;
```

```
create public synonym prepare_connection for prepare_connection; grant execute on prepare_connection to public; ""
```

10.1.4 Zaxxer

TODO how to call this procedure in the connection pool

10.2 ## DBMS_SESSION

10.3 Identifier

SET_IDENTIFIER and CLEAR_IDENTIFIER procedures to allow the real user to be associated with a session, regardless of what database user was being used for the connection.

10.4 Metrics

```
try { String    e2eMetrics[]    =    new    String[OracleConnection.END_TO_END_STATE_INDEX_MAX];
    e2eMetrics[OracleConnection.END_TO_END_ACTION_INDEX] = null; e2eMetrics[OracleConnection.END_TO_END_MOD
    = null; e2eMetrics[OracleConnection.END_TO_END_CLIENTID_INDEX] = null; ((OracleConnection)
    conn).setEndToEndMetrics(e2eMetrics, Short.MIN_VALUE);
```

```
} catch (SQLException sqle) { // Do something...
```

```
}
```

0 - No trace. Like switching sql_trace off. 2 - The equivalent of regular sql_trace. 4 - The same as 2, but with the addition of bind variable values. 8 - The same as 2, but with the addition of wait events. 12 - The same as 2, but with both bind variable values and wait events.

Monitoring long running <https://oracle-base.com/articles/11g/real-time-sql-monitoring-11gr1>

Database Objects

Entity Relationship Diagram [logger_tables.png](#)

Table DDL

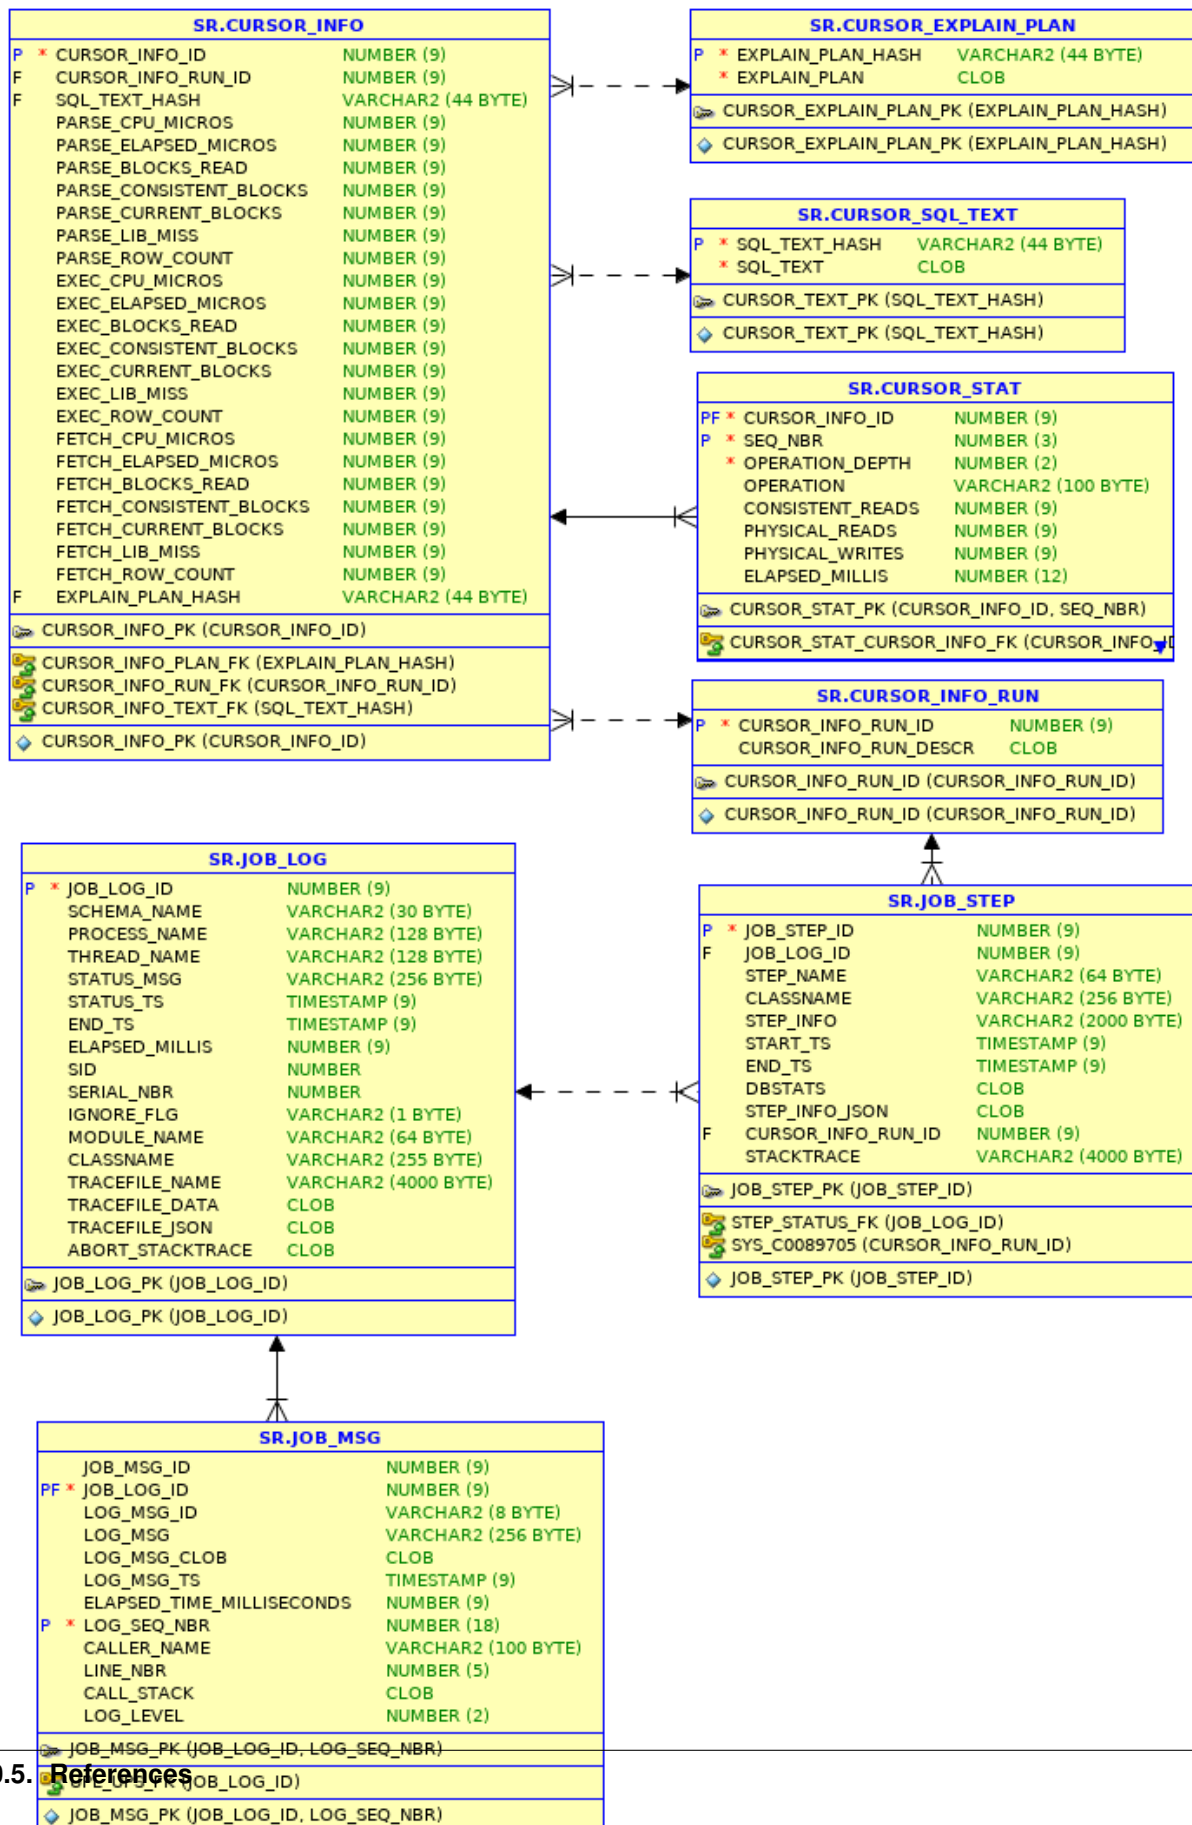
Oracle

10.5 References

!![javadoc] ('../target/site/apidocs/index.html')

https://oracle-base.com/articles/misc/dbms_session

<https://oracle-base.com/articles/misc/sql-trace-10046-trcsess-and-tkprof>



SPRING DEVELOPERS

Oracle tracing is a powerful tool that logs detailed information about all calls to the Oracle database.

In order to use this :

- one must turn on tracing for the current connection
- set the log file
- stop tracing
- call a service to store the trace
- store the raw trace file
- analyze the trace file
- store the analyzed trace file

Logs messages using `utl_file` to a directory on the database server specified.

First the database directory is created and oracle is granted permission to read and write it, then the ddl “create directory...” and “grant read, write on directory...”

```
set serveroutput on
set echo on
create or replace procedure log_to_file_only is
    long_msg clob := 'this is an absurdly long message, ' ||
        ' interesting stuff to say so I will just write meaningless ' ||
        ' stuff for a little while. ' ||
        ' The quick brown fox jumped over the lazy dog. ';

    my_log_file_name varchar(4096);
begin
    my_log_file_name := pllogger.open_log_file('log_to_file_only.text');
    pllogger.set_filter_level(9); -- all messages should go to log file
    pllogger.info('anonymous', $$PLSQL_LINE, 'begin loop');
    pllogger.info($$PLSQL_UNIT, $$PLSQL_LINE, long_msg);
    for i in 1..3
    loop
        pllogger.fine($$PLSQL_UNIT, $$PLSQL_LINE, 'i is ' || to_char(i));
    end loop;
    pllogger.close_log_file();
exception when others then
    -- a severe condition is not necessarily fatal
    pllogger.severe($$PLSQL_UNIT, $$PLSQL_LINE, sqlerrm);
    pllogger.close_log_file();
    raise;
```

```
end;
/
show errors

exec log_to_file_only();

"log_level","job_log_id","job_msg_id","line_number","timestamp","log_msg","caller_name","call_stack"
4,,,17,"2019-10-26T17:19:52.885607","begin loop","anonymous",""
4,,,18,"2019-10-26T17:19:52.886020","this is an absurdly long message, exceeding the length of the l
7,,,22,"2019-10-26T17:19:52.886197","i is 1","LOG_TO_FILE_ONLY",""
7,,,22,"2019-10-26T17:19:52.886357","i is 2","LOG_TO_FILE_ONLY",""
7,,,22,"2019-10-26T17:19:52.886502","i is 3","LOG_TO_FILE_ONLY",""
```

TODO

Tracing should do the following

- Begin with any transaction as annotated by @Transactional

INSTALL

```
cd src/main/resources/ddl/oracle
```

```
sqlplus $ORACLE_UID @ pllogger.pkgs.sr.sql
```

```
sqlplus $ORACLE_UID @ pllogger.pkgb.sr.sql
```

create directory job_msg_dir as '/common/scratch/ut_process_log_dir'; grant write on directory to sr;

should be granted by user, not by role.

- Configuring to use your database
- Example schema

TRACE FILE FIELDS

14.1 Trace Record Fields

code-block:

		Cursor Operation								
	Type	Parsing	ParseError	Parse	EXec	Fetch	Stat	Lobread	Lobpgsize	Close
#	cursorNumbe		X			X	X			
ad	sgaAddress	X								
bytes	bytes							X	X	
c	CpuMicroSec			X	X	X		X	X	X
card	cardinality									
cnt							O			
cost	cost (optim									
cr	consistentR			X	X	X		X	X	
cu	currentMode			X	X	X		X	X	
dep	depth	X	X	X	X	X				X
e	elapsedMicr			X	X	X		X	X	X
err	oracleError		X							
hv	sqlHashValu	X								
id							X			
len	sqlTeXtLeng	X	X							
lid		X	X							
mis	libraryCach			X	X	X				
obj	objectNumbe						X			
oct	oracleComma	X	X							
og	optimizerGo			X	X	X				
op	operation						X			
p	physicalBlo			X	X	X		X	X	
pid	processId						X			
plh				X	X	X				
pos	position (o						X			
pw	physicalWri						X			
r	rowCount			X	X	X				
size							O			
sqlid	sqlId	X								
str								X		
tim				X	X	X		X	X	X
time										
type										
uid		X	X							

TODO

- security can't specify file name
- need an agent to get the log files for remote users
- TODO escape double quotes in text fields
- check for anomolous run-times using condition identification
- plot runtimes
- TODO describe microservices, multiple connections, tying them all together
- TODO describe using with spring